

- Name: Gokulakrishnan B
- Roll No: DA24M007
- Department: Mtech DSAI
- Assignment No: 8


## Importing libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.colors as colors
from matplotlib.colors import ListedColormap
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.base import clone
```

## Creating dataset (given in question)

```
In [2]: from sklearn.datasets import make_moons, make_circles
from sklearn.model_selection import train_test_split

X, y = make_circles(n_samples=500, noise=0.1, random_state=42, factor=0.2)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
y_train = np.where(y_train == 0, -1, y_train)
y_test = np.where(y_test == 0, -1, y_test)
plt.scatter(X[:,0], X[:,1], c=y, cmap=colors.ListedColormap(["blue", "red"]))
plt.axis('equal')
plt.show()
```



```
In [3]: pd.Series(y).value_counts()
```

```
Out[3]: 1    250
0     250
Name: count, dtype: int64
```

The dataset is balanced.

## Implementing AdaBoost Algorithm from scratch

```
In [4]: class Adaboost:
def __init__(self, base_learner, n_learners, eta=0.5, **learner_params):
self.base_learner = base_learner
self.n_learners = n_learners
self.eta = eta
self.learners = []
self.alphas = []
self.errors = []
self.learner_params = learner_params
```

```

def fit(self, X, y):
    n_samples, n_features = X.shape
    w = np.ones(n_samples) / n_samples

    for _ in range(self.n_learners):
        learner = self.base_learner(**self.learner_params)
        learner.fit(X, y, sample_weight=w)
        y_pred = learner.predict(X)
        error = np.sum(w[y != y_pred]) / np.sum(w)

        alpha = self.eta * 0.5 * np.log((1 - error) / (error + 1e-10))

        w = w * np.exp(-alpha * y_pred * y)
        w = w / np.sum(w)

        self.alphas.append(alpha)
        self.learners.append(learner)
        self.errors.append(error)

def predict(self, X):
    predictions = np.zeros(X.shape[0])

    for alpha, learner in zip(self.alphas, self.learners):
        predictions += alpha * learner.predict(X)

    predictions = np.sign(predictions)

    return predictions

def plot_decision_boundaries(self, X, y):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                          np.arange(y_min, y_max, 0.02))

    cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
    cmap_bold = ['#FF0000', '#0000FF']

    plt.figure(figsize=(8, 6))

    for i, learner in enumerate(self.learners):
        Z = learner.predict(np.c_[xx.ravel(), yy.ravel()])
        Z = Z.reshape(xx.shape)

        plt.contour(xx, yy, Z, alpha=0.3, linewidths=1)

    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=ListedColormap(cmap_bold), edgecolor='k', s=30)

    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.title('Combined Decision Boundaries of Weak Learners')
    plt.xlabel('X1')
    plt.ylabel('X2')

    plt.show()

def plot_final_decision_boundary(self, X, y):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                          np.arange(y_min, y_max, 0.02))

    Z = self.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
    cmap_bold = ['#FF0000', '#0000FF']

    plt.contourf(xx, yy, Z, cmap=cmap_light)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=ListedColormap(cmap_bold), edgecolor='k', s=20)
    plt.title('Final AdaBoost Decision Boundary')
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.tight_layout()
    plt.show()

```

## Hyperparameter tuning on Adaboost algorithm

```

In [5]: from sklearn.metrics import accuracy_score
# Model - Decision stump

eta_values = np.arange(0,1,0.1)
n_learners_values=np.arange(1,10,2)
baseClassifier = DecisionTreeClassifier

best_eta=None
best_n_learner = None
max_score = float('-inf')

for eta in eta_values:
    for n_learners in n_learners_values:
        model = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,max_depth=1)
        model.fit(X_train,y_train)
        y_pred = model.predict(X_test)
        score = accuracy_score(y_test,y_pred)

        if score > max_score:
            max_score = score
            best_eta=eta
            best_n_learner = n_learners


print(f"Best eta value : {best_eta}")
print(f"Best n_learners value : {best_n_learner}")
print(f"Accuracy score on best parameters: {max_score}")


ada = Adaboost(baseClassifier,best_n_learner,best_eta)
ada.fit(X_train,y_train)
ada.plot_decision_boundaries(X,y)
ada.plot_final_decision_boundary(X,y)

print(classification_report(y_test,ada.predict(X_test)))

```

Best eta value : 0.8  
 Best n\_learners value : 9  
 Accuracy score on best parameters: 0.904





	precision	recall	f1-score	support
-1	0.97	1.00	0.98	61
1	1.00	0.97	0.98	64
accuracy			0.98	125
macro avg	0.98	0.98	0.98	125
weighted avg	0.98	0.98	0.98	125

Here, we used a decision stump (a decision tree with just one node) as the weak learner. As shown in the images above, the boosted model is able to recognize the pattern in the data more effectively than the individual weak learners, although it doesn't capture the relationship perfectly. To enhance its performance, we can experiment with different models as weak learners.

```
In [6]: from sklearn.metrics import accuracy_score
# Model - Logistic Regression

eta_values = np.arange(0,1,0.1)
n_learners_values=np.arange(1,10,2)
param_grid_log_penalty = ['l1','l2']
param_grid_log_c = [1,3,5]
baseClassifier = LogisticRegression

best_eta=None
best_n_learner = None
max_score = float('-inf')
best_penalty=None
best_c=None

for eta in eta_values:
    for n_learners in n_learners_values:
        for penalty in param_grid_log_penalty:
            for c in param_grid_log_c:
                model = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,solver='liblinear', random_state=42,
                                model.fit(X_train,y_train)
                                y_pred = model.predict(X_test)
                                score = accuracy_score(y_test,y_pred)



                if score > max_score:
                    max_score = score
                    best_eta=eta
                    best_n_learner = n_learners
                    best_penalty=penalty
                    best_c=c

print(f"Best eta value : {best_eta}")
print(f"Best n_learners value : {best_n_learner}")
print(f"Accuracy score on best parameters: {max_score}")

ada = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,solver='liblinear', random_state=42,max_iter=10000,pen
ada.fit(X_train,y_train)
ada.plot_decision_boundaries(X,y)
ada.plot_final_decision_boundary(X,y)

print(classification_report(y_test,ada.predict(X_test)))
```

Best eta value : 0.8  
 Best n\_learners value : 3  
 Accuracy score on best parameters: 0.736



	precision	recall	f1-score	support
-1	0.85	0.46	0.60	61
1	0.64	0.92	0.76	64
accuracy			0.70	125
macro avg	0.74	0.69	0.68	125
weighted avg	0.74	0.70	0.68	125

```
In [7]: from sklearn.metrics import accuracy_score
# Model - Decision Tree

eta_values = np.arange(0,1,0.1)
n_learners_values=np.arange(1,10,2)
param_grid_dt = [3,5,7,9] # min sample split
baseClassifier = DecisionTreeClassifier

best_eta=None
best_n_learner = None
max_score = float('-inf')
best_dt = None

for eta in eta_values:
    for n_learners in n_learners_values:
        for dt in param_grid_dt:
            model = AdaBoost(baseClassifier,n_learners=n_learners, eta=eta,min_samples_split=dt,max_depth=3)
```

```

model.fit(X_train,y_train)
y_pred = model.predict(X_test)
score = accuracy_score(y_test,y_pred)

if score > max_score:
    max_score = score
    best_eta=eta
    best_n_learner = n_learners
    best_dt = dt

print(f"Best eta value : {best_eta}")
print(f"Best n_learners value : {best_n_learner}")
print(f"Accuracy score on best parameters: {max_score}")

ada = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,min_samples_split=best_dt,max_depth=3)
ada.fit(X_train,y_train)
ada.plot_decision_boundaries(X,y)
ada.plot_final_decision_boundary(X,y)

print(classification_report(y_test,ada.predict(X_test)))


```

Best eta value : 0.1


Best n\_learners value : 3

Accuracy score on best parameters: 1.0

Combined Decision Boundaries of Weak Learners



Final AdaBoost Decision Boundary





	precision	recall	f1-score	support
-1	0.98	1.00	0.99	61
1	1.00	0.98	0.99	64
accuracy			0.99	125
macro avg	0.99	0.99	0.99	125
weighted avg	0.99	0.99	0.99	125

```
In [8]: from sklearn.metrics import accuracy_score
# Model - Linear SVM

eta_values = np.arange(0,1,0.1)
n_learners_values=np.arange(1,10,2)
param_grid_svc_c = [3,5,7,9]
baseClassifier = SVC

best_eta=None
best_n_learner = None
max_score = float('-inf')
best_c = None

for eta in eta_values:
    for n_learners in n_learners_values:
        for c in param_grid_svc_c:
            model = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,C=c,kernel='linear')
            model.fit(X_train,y_train)
            y_pred = model.predict(X_test)
            score = accuracy_score(y_test,y_pred)


            if score > max_score:
                max_score = score
                best_eta=eta
                best_n_learner = n_learners
                best_c = c


print(f"Best eta value : {best_eta}")
print(f"Best n_learners value : {best_n_learner}")
print(f"Accuracy score on best parameters: {max_score}")

ada = Adaboost(baseClassifier,n_learners=n_learners, eta=eta,C=best_c,kernel='linear')
ada.fit(X_train,y_train)
ada.plot_decision_boundaries(X,y)
ada.plot_final_decision_boundary(X,y)

print(classification_report(y_test,ada.predict(X_test)))
```

Best eta value : 0.1  
 Best n\_learners value : 1  
 Accuracy score on best parameters: 0.488





	precision	recall	f1-score	support
-1	0.49	1.00	0.66	61
1	0.00	0.00	0.00	64
accuracy			0.49	125
macro avg	0.24	0.50	0.33	125
weighted avg	0.24	0.49	0.32	125

```

/home/gokul/miniconda3/envs/dsai/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/gokul/miniconda3/envs/dsai/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/home/gokul/miniconda3/envs/dsai/lib/python3.10/site-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

## Incorporating Various weak models into a single Adaboost model

```

In [9]: from sklearn.linear_model import LogisticRegression
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.svm import SVC
        from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
        from sklearn.base import clone

```

```

In [10]: from sklearn.base import BaseEstimator, ClassifierMixin

        class NewAdaboost(BaseEstimator, ClassifierMixin):
            def __init__(self, base_learners, n_learners=50, eta=0.5):
                self.base_learners = base_learners
                self.n_learners = n_learners
                self.eta = eta
                self.learners = []
                self.alphas = []
                self.errors=[]

            def fit(self, X, y):
                n_samples = X.shape[0]
                w = np.ones(n_samples) / n_samples

                for i in range(self.n_learners):
                    learner = clone(self.base_learners[i % len(self.base_learners)])
                    learner.fit(X, y, sample_weight=w)
                    y_pred = learner.predict(X)

                    error = np.sum(w * (y_pred != y)) / np.sum(w)

                    alpha = self.eta * np.log((1 - error) / error)

                    w = w * np.exp(-alpha * y * y_pred)
                    w = w / np.sum(w)

```



```

        self.learners.append(learner)
        self.alphas.append(alpha)
        self.errors.append(error)

def predict(self, X):
    final_predictions = np.zeros(X.shape[0])
    for alpha, learner in zip(self.alphas, self.learners):
        final_predictions += alpha * learner.predict(X)

    return np.sign(final_predictions)

def plot_decision_boundaries(self, X, y):
    # Define grid for plotting
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                          np.arange(y_min, y_max, 0.02))

    # Set up the plot
    cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
    cmap_bold = ['#FF0000', '#0000FF']

    plt.figure(figsize=(8, 6))

    # Loop over each weak learner and plot its decision boundary
    for i, learner in enumerate(self.learners):
        Z = learner.predict(np.c_[xx.ravel(), yy.ravel()])
        Z = Z.reshape(xx.shape)

        # Plot decision boundary (contour lines)
        plt.contour(xx, yy, Z, alpha=0.3, linewidths=1) # Alpha and linewidth adjust the visibility of the con

    # Plot the data points
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=ListedColormap(cmap_bold), edgecolor='k', s=30)

    # Labels and limits
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.title('Combined Decision Boundaries of Weak Learners')
    plt.xlabel('X1')
    plt.ylabel('X2')

    plt.show()

def plot_final_decision_boundary(self, X, y):
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.02),
                          np.arange(y_min, y_max, 0.02))

    Z = self.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    cmap_light = ListedColormap(['#FFAAAA', '#AAAAFF'])
    cmap_bold = ['#FF0000', '#0000FF']

    plt.contourf(xx, yy, Z, cmap=cmap_light)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=ListedColormap(cmap_bold), edgecolor='k', s=20)
    plt.title('Final AdaBoost Decision Boundary')
    plt.xlim(xx.min(), xx.max())
    plt.ylim(yy.min(), yy.max())
    plt.tight_layout()
    plt.show()

```

```

In [11]: weak_learners = [
    DecisionTreeClassifier(max_depth=1, random_state=42),
    DecisionTreeClassifier(max_depth=3, random_state=42),
    LogisticRegression(solver='liblinear', random_state=42, max_iter=10000),
    SVC(kernel='linear', random_state=42)
]

```

## Hyperparameter tuning on various parameters of weak learners.

```

In [12]: eta_values = np.arange(0,1,0.3)
    n_learners_values=np.arange(1,20,5)

    param_grid_dt = [3,5] # min samples split
    param_grid_log_penalty = ['l1','l2'] # penalty
    param_grid_log_c = [1,3,5] # c value in Logistic Regression
    param_grid_svc_c = [1,3,5] # c value in SVC

```

In [13]: `from sklearn.metrics import accuracy_score`

```
best_eta=None
best_n_learner = None
best_dt1 = None
best_dt2= None
best_log_penalty = None
best_log_c = None
best_svc_c = None
max_score = float('-inf')

for eta in eta_values:
    for n_learners in n_learners_values:
        for dt1 in param_grid_dt:
            for dt2 in param_grid_dt:
                for log_penalty in param_grid_log_penalty:
                    for log_c in param_grid_log_c:
                        for svc_c in param_grid_svc_c:

                            weak_learners = [
                                DecisionTreeClassifier(max_depth=1, random_state=42,min_samples_split=dt1),
                                DecisionTreeClassifier(max_depth=3, random_state=42,min_samples_split=dt2),
                                LogisticRegression(solver='liblinear', random_state=42,max_iter=10000, penalty=log_
                                SVC(kernel='linear', random_state=42,C=svc_c)
                            ]

                            model = NewAdaboost(base_learners=weak_learners,n_learners=n_learners, eta=eta)
                            model.fit(X_train,y_train)
                            y_pred = model.predict(X_test)
                            score = accuracy_score(y_test,y_pred)

                            if score > max_score:
                                max_score = score
                                best_eta=eta
                                best_n_learner = n_learners
                                best_dt1=dt1
                                best_dt2=dt2
                                best_log_penalty=log_penalty
                                best_log_c = log_c
                                best_svc_c=svc_c



print(f"Best eta value : {best_eta}")
print(f"Best n_learners value : {best_n_learner}")
print(f"Best dt1 value : {best_dt1}")
print(f"Best dt2 value : {best_dt2}")
print(f"Best penalty for log reg : {best_log_penalty}")
print(f"Best c value for log reg : {best_log_c}")
print(f"Best c value for SVC : {best_svc_c}")
print(f"Accuracy score on best parameters: {max_score}")
```

Best eta value : 0.6  
Best n\_learners value : 11  
Best dt1 value : 3  
Best dt2 value : 3  
Best penalty for log reg : 11  
Best c value for log reg : 5  
Best c value for SVC : 3  
Accuracy score on best parameters: 1.0

In [14]: `weak_learners = [`

```
                                DecisionTreeClassifier(max_depth=1, random_state=42,min_samples_split=best_dt1),
                                DecisionTreeClassifier(max_depth=3, random_state=42,min_samples_split=best_dt2),
                                LogisticRegression(solver='liblinear', random_state=42,max_iter=10000, penalty=best
                                SVC(kernel='linear', random_state=42,C=best_svc_c)
                                ]

ada = NewAdaboost(weak_learners,best_n_learner,best_eta)
ada.fit(X_train,y_train)
ada.plot_decision_boundaries(X,y)
ada.plot_final_decision_boundary(X,y)
```



```
In [15]: print(classification_report(y_test, ada.predict(X_test)))
```

	precision	recall	f1-score	support
-1	1.00	1.00	1.00	61
1	1.00	1.00	1.00	64
accuracy			1.00	125
macro avg	1.00	1.00	1.00	125
weighted avg	1.00	1.00	1.00	125

Here, we observe a similar outcome to using only decision stumps as weak learners, but this approach has outperformed using a single type of weak learner. By utilizing multiple weak learners, we achieved 100% accuracy.

```
In [16]: print('Alpha Values')
print(*[float(i) for i in ada.alphas])
print('')
print('Error Values')
print(*[float(i) for i in ada.errors])
```

Alpha Values  
0.48944970159136664 1.8448097080522923 0.9414683502754397 0.9078059743192963 0.5095112286745234 1.8806960805637407 0.8156227255665596 0.3521439023874767 0.8155256001830995 1.9069936651164854 0.3117266741365387

Error Values  
0.30666666666666664 0.044163726705591294 0.17234274654376913 0.1804931459753935 0.2996037703412075 0.04170664543715051 0.20434251625702005 0.3573449603046424 0.20436883633488606 0.03998967379547085 0.372958761069171

Here, we can observe that, the alpha values are lower for weak learners that have high errors, and vice versa.

In [ ]: